

Introducing Pandas Objects

At a very **basic** level, **Pandas objects** can be thought of as **enhanced versions of NumPy structured arrays**.

In which the **rows and columns** are **identified with labels rather than simple integer indices**.

As we will see during the course of this chapter, **Pandas provides** a host of **useful tools, methods, and functionality** on **top of the basic data structures**.

But nearly **everything** that follows will require an **understanding of what these structures are**.

Thus, before we go any further, let's take a look at these** three fundamental Pandas data structures:** the **Series**, **DataFrame**, and **Index**.

We will start our code sessions with the **standard NumPy and Pandas imports**:

```
In [1]: import numpy as np  
import pandas as pd
```

The Pandas Series Object

A Pandas **Series** is a **one-dimensional array of indexed data**.

It can be **created from a list or array** as follows:

```
In [ ]: data = pd.Series([0.25, 0.5, 0.75, 1.0])  
data
```

```
Out[ ]: 0    0.25  
1    0.50  
2    0.75  
3    1.00  
dtype: float64
```

The **Series** combines a sequence of values with an explicit **sequence of indices**, which we can **access** with the **values** and **index** attributes.

The **values** are simply a familiar NumPy array:

```
In [ ]: data.values
```

```
Out[ ]: array([0.25, 0.5 , 0.75, 1.  ])
```

The **index** is an **array-like object** of type **pd.Index**, which we'll discuss in more detail momentarily:

```
In [ ]: data.index
```

```
Out[ ]: RangeIndex(start=0, stop=4, step=1)
```

Like with a NumPy array, **data can be accessed** by the associated **index** via the familiar Python **square-bracket notation**:

```
In [ ]: data[1]
```

```
Out[ ]: 0.5
```

```
In [ ]: data[1:3]
```

```
Out[ ]: 1    0.50  
       2    0.75  
       dtype: float64
```

As we will see, though, the Pandas **Series** is **much more general** and **flexible** than the **one-dimensional NumPy array** that it emulates.

Series as Generalized NumPy Array

From what we've seen so far, the **Series** object **may appear** to be basically **interchangeable** with a **one-dimensional NumPy array**.


```
Out[ ]:  a    0.25  
        b    0.50  
        c    0.75  
        d    1.00  
        dtype: float64
```

And the **item access works as expected:**

```
In [ ]: data['b']
```

```
Out[ ]: 0.5
```

We can even use **noncontiguous** or **nonsequential indices**:

```
In [ ]: data = pd.Series([0.25, 0.5, 0.75, 1.0],  
                        index=[2, 5, 3, 7])  
data
```

```
Out[ ]: 2    0.25  
        5    0.50  
        3    0.75  
        7    1.00  
        dtype: float64
```

```
In [ ]: data[5]
```

```
Out[ ]: 0.5
```

Series as Specialized Dictionary

In this way, you can think of a Pandas **Series** a bit like a **specialization of a Python dictionary**.

A **dictionary** is a structure that **maps arbitrary keys to a set of arbitrary values**, and a **Series** is a structure that **maps typed keys to a set of typed values**.

This **typing is important**:

Just as the **type-specific** compiled code behind a NumPy array makes it more **efficient** than a Python list for certain operations,

the **type information** of a Pandas **Series** makes it **more efficient than Python dictionaries** for certain operations.

The **Series -as-dictionary analogy** can be made even **more clear** by **constructing a Series object directly from a Python dictionary**.

Here the **five most populous US states** according to the 2020 census:

```
In [ ]: population_dict = {'California': 39538223, 'Texas': 29145505,  
                           'Florida': 21538187, 'New York': 20201249,  
                           'Pennsylvania': 13002700}  
  
population = pd.Series(population_dict)  
population
```



```
Out[ ]: California    39538223
        Texas         29145505
        Florida       21538187
        New York      20201249
        Pennsylvania   13002700
        dtype: int64
```

From here, **typical dictionary-style item access can be performed:**

```
In [ ]: population['California']
```

```
Out[ ]: 39538223
```

Unlike a dictionary, though, the **Series** also **supports** array-style operations such as **slicing**:

```
In [ ]: population['California':'Florida']
```

```
Out[ ]: California    39538223  
        Texas        29145505  
        Florida      21538187  
        dtype: int64
```

We'll discuss some of the quirks of **Pandas indexing and slicing** in **Data Indexing and Selection**.

Constructing Series Objects

We've already seen a **few ways of constructing a Pandas Series from scratch**. All of them are **some version of the following**:

```
pd.Series(data, index=index)
```

where **index** is an **optional** argument, and **data** can be **one of many entities**.

For example, data can be a list or NumPy array, in which case **index** defaults to an integer sequence:

```
In [ ]: pd.Series([2, 4, 6])
```

```
Out[ ]: 0    2  
        1    4  
        2    6  
        dtype: int64
```

Or **data** can be a **scalar**, which is **repeated** to fill the **specified index**:

```
In [ ]: pd.Series(5, index=[100, 200, 300])
```

```
Out[ ]: 100    5  
        200    5  
        300    5  
        dtype: int64
```

Or it can be a **dictionary**, in which case **index** defaults to the **dictionary keys**:

```
In [ ]: pd.Series({2: 'a', 1: 'b', 3: 'c'})
```

```
Out[ ]: 2    a
        1    b
        3    c
        dtype: object
```

In **each case**, the **index can be explicitly set to control the order** or the **subset of keys** used:

```
In [ ]: pd.Series({2:'a', 1:'b', 3:'c'}, index=[1, 2])
```

```
Out[ ]: 1    b
        2    a
        dtype: object
```

The Pandas DataFrame Object

The next **fundamental structure** in Pandas is the **DataFrame**.

Like the **Series** object discussed in the previous section, the **DataFrame** can be thought of either as a **generalization of a NumPy array**, or as a **specialization of a Python dictionary**.

We'll now take a look at each of these perspectives.

DataFrame as Generalized NumPy Array

If a **Series** is an analog of a one-dimensional array with explicit indices, a **DataFrame** is an analog of a two-dimensional array with explicit row and column indices.

Just as you might think of a **two-dimensional array** as an **ordered sequence of aligned one-dimensional columns**, you can think of a **DataFrame** as a sequence of aligned **Series objects**.

Here, by "**aligned**" we mean that they share the same index.

To demonstrate this, let's first **construct a new Series** listing **the area of each of the five states** discussed in the previous section (in square kilometers):

```
In [ ]: area_dict = {'California': 423967, 'Texas': 695662, 'Florida': 170312,
                    'New York': 141297, 'Pennsylvania': 119280}
area = pd.Series(area_dict)
area
```

```
Out[ ]: California    423967
Texas              695662
Florida           170312
New York          141297
Pennsylvania      119280
dtype: int64
```

Now that we have this **along with the population Series** from before,

we can **use a dictionary** to construct a **single two-dimensional object containing this information**:

```
In [ ]: states = pd.DataFrame({'population': population,
                              'area': area})
states
```

Out[]:

	population	area
California	39538223	423967
Texas	29145505	695662
Florida	21538187	170312
New York	20201249	141297
Pennsylvania	13002700	119280

Like the `Series` object, the `DataFrame` has an `index` **attribute** that **gives access to the index labels**:

```
In [ ]: states.index
```

```
Out[ ]: Index(['California', 'Texas', 'Florida', 'New York', 'Pennsylvania'], dtype='object')
```

Additionally, the `DataFrame` has a `columns` **attribute**, which is an `Index` **object holding the column labels**:

```
In [ ]: states.columns
```

```
Out[ ]: Index(['population', 'area'], dtype='object')
```

Thus the **DataFrame** can be thought of as a **generalization of a two-dimensional NumPy array**,

where **both the rows and columns** have a **generalized index for accessing the data**.

DataFrame as Specialized Dictionary

Similarly, we can also think of a **DataFrame** as a **specialization of a dictionary**.

Where a **dictionary maps a key to a value**, a **DataFrame** maps a **column name to a Series** of column data.

For example, asking for the **'area'** attribute returns the **Series object** containing the areas we saw earlier:


```
In [ ]: states['area']
```

```
Out[ ]: California      423967  
Texas      695662  
Florida    170312  
New York   141297  
Pennsylvania 119280  
Name: area, dtype: int64
```

Notice the **potential point of confusion** here:

in a **two-dimensional NumPy array**, `data[0]` will return the **first row**.

For a **DataFrame**, `data['col0']` will return the first column.

Because of this, it is probably **better to think about DataFrame s** as **generalized dictionaries rather than generalized arrays**,

though **both** ways of looking at the situation **can be useful**.

We'll explore **more flexible means of indexing** `DataFrame` s in **Data Indexing and Selection**.

Constructing DataFrame Objects

A Pandas `DataFrame` can be constructed in a variety of ways.

Here we'll explore **several examples**.

From a single Series object

A `DataFrame` is a collection of `Series` objects,

and a **single-column** `DataFrame` can be **constructed from a single** `Series` :

```
In [ ]: pd.DataFrame(population, columns=['population'])
```

Out[]:

	population
California	39538223
Texas	29145505
Florida	21538187
New York	20201249
Pennsylvania	13002700

From a list of dicts

Any list of dictionaries can be made into a `DataFrame` .

We'll use a simple **list comprehension to create some data:**

```
In [2]: data = [{'a': i, 'b': 2 * i}
            for i in range(3)]
data
```

```
Out[2]:  [{'a': 0, 'b': 0}, {'a': 1, 'b': 2}, {'a': 2, 'b': 4}]
```

```
In [3]: pd.DataFrame(data)
```

```
Out[3]:
```

	a	b
0	0	0
1	1	2
2	2	4

Even if **some keys** in the dictionary are **missing**,

Pandas will **fill** them in with **NaN values** (i.e., "Not a Number"; see **Handling Missing Data**):

```
In [ ]: pd.DataFrame([{'a': 1, 'b': 2}, {'b': 3, 'c': 4}])
```

Out[]:

	a	b	c
0	1.0	2	NaN
1	NaN	3	4.0

From a dictionary of Series objects

As we saw before, a **DataFrame** can be constructed from a **dictionary of Series objects** as well:

```
In [ ]: pd.DataFrame({'population': population,  
                      'area': area})
```

Out[]:

	population	area
California	39538223	423967
Texas	29145505	695662
Florida	21538187	170312
New York	20201249	141297
Pennsylvania	13002700	119280

From a two-dimensional NumPy array

Given a **two-dimensional array of data**,

we can create a **DataFrame** with any specified column and index names.

If omitted, an integer index will be used for each:

```
In [4]: np.random.rand(3, 2)
```

```
Out[4]: array([[0.89150021, 0.52792497],  
               [0.55672672, 0.19928865],  
               [0.61379592, 0.72308653]])
```

```
In [ ]: pd.DataFrame(np.random.rand(3, 2),  
                      columns=['foo', 'bar'],  
                      index=['a', 'b', 'c'])
```

```
Out[ ]:
```

	foo	bar
a	0.471098	0.317396
b	0.614766	0.305971
c	0.533596	0.512377

From a NumPy structured array

We covered structured arrays in **Structured Data: NumPy's Structured Arrays**.

A Pandas `DataFrame` operates much like a structured array, and **can be created directly from one:**

```
In [ ]: A = np.zeros(3, dtype=[('A', 'i8'), ('B', 'f8')])  
A
```

```
Out[ ]: array([(0, 0.), (0, 0.), (0, 0.)], dtype=[('A', '<i8'),  
          ('B', '<f8')])
```

```
In [ ]: pd.DataFrame(A)
```


Out[]:

	A	B
0	0	0.0
1	0	0.0
2	0	0.0

The Pandas Index Object

As you've seen, the **Series** and **DataFrame** objects both **contain an explicit index** that lets you reference and modify data.

This **Index** object is an **interesting structure** in itself,

and it can be **thought of** either as an **immutable array** or as an **ordered set** (technically a **multiset**, as **Index** objects may contain repeated values).

Those **views** have some **interesting consequences** in terms of the **operations available on Index objects**.

As a simple **example**, let's **construct an Index from a list of integers**:

```
In [ ]: ind = pd.Index([2, 3, 5, 7, 11])  
ind
```

```
Out[ ]: Int64Index([2, 3, 5, 7, 11], dtype='int64')
```

Index as Immutable Array

The **Index** in many ways operates like an array.

For **example**, we can use **standard Python indexing notation to retrieve values or slices**:

```
In [ ]: ind[1]
```

```
Out[ ]: 3
```

```
In [ ]: ind[:,2]
```

```
Out[ ]: Int64Index([2, 5, 11], dtype='int64')
```

Index objects also have **many of the attributes familiar from NumPy arrays**:

```
In [ ]: print(ind.size, ind.shape, ind.ndim, ind.dtype)
```

```
5 (5,) 1 int64
```

One **difference between Index objects and NumPy arrays** is that the **indices are immutable** — that is, they **cannot be modified via the normal means**:

```
In [ ]: ind[1] = 0
```

```

-----
-----
TypeError                                Traceback (most recent call last)
/var/folders/xc/sptt9bk14s34rgxt7453p03r0000gp/T/ipykernel_83282/393126374.py in <module>
----> 1 ind[1] = 0

~/local/share/venv/python-data-science-handbook-2e-ukwqDTB/lib/python3.9/site-packages/pandas/core/indexes/base.py in __setitem__(self, key, value)
    4583     @final
    4584     def __setitem__(self, key, value):
-> 4585         raise TypeError("Index does not support mutable operations")
    4586
    4587     def __getitem__(self, key):

TypeError: Index does not support mutable operations

```

This **immutability** makes it safer to share indices between multiple `DataFrame`s and arrays,

without the potential for **side effects** from **inadvertent index modification**.

Index as Ordered Set

Pandas objects are designed to **facilitate operations** such as **joins across datasets**, which depend on many aspects of set arithmetic.

The **Index object** follows many of the **conventions** used by **Python's built-in set** data structure,

so that **unions, intersections, differences, and other combinations** can be computed in a familiar way:

```
In [ ]: indA = pd.Index([1, 3, 5, 7, 9])  
        indB = pd.Index([2, 3, 5, 7, 11])
```

```
In [ ]: indA.intersection(indB)
```

```
Out[ ]: Int64Index([3, 5, 7], dtype='int64')
```

```
In [ ]: indA.union(indB)
```

```
Out[ ]: Int64Index([1, 2, 3, 5, 7, 9, 11], dtype='int64')
```

```
In [ ]: indA.symmetric_difference(indB)
```

```
Out[ ]: Int64Index([1, 2, 9, 11], dtype='int64')
```