

Aggregations: min, max, and Everything in Between

A **first step** in **exploring any dataset** is often to compute various **summary statistics**.

Perhaps the **most common summary statistics** are the **mean** and standard **deviation**.

Which allow you to **summarize** the "**typical**" **values** in a dataset, but other aggregations are useful as well (the sum, product, median, minimum and maximum, quantiles, etc.).

NumPy has **fast** built-in **aggregation** functions for working on arrays.

Summing the Values in an Array

As a **quick example**, consider computing **the sum of all values** in an array.

Python itself can do this using the **built-in** `sum` **function**:

```
In [7]: import numpy as np  
rng = np.random.default_rng()
```

```
In [ ]: L = rng.random(100)  
sum(L)
```

```
Out[ ]: 52.76825337322368
```

The **syntax** is quite **similar** to that of **NumPy's** `sum` **function**, and the result is the same in the simplest case:

```
In [ ]: np.sum(L)
```

```
Out[ ]: 52.76825337322366
```

However, because it executes the operation in **compiled** code, **NumPy's version** of the operation is computed **much more quickly**:

```
In [ ]: big_array = rng.random(1000000)
        %timeit sum(big_array)
        %timeit np.sum(big_array)
```

89.9 ms \pm 233 μ s per loop (mean \pm std. dev. of 7 runs, 10 loops each)

521 μ s \pm 8.37 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

Be careful, though:

- The `sum` function and the `np.sum` function are **not identical**, which can sometimes lead to confusion!
- In particular, their **optional arguments** have **different meanings** (`sum(x, 1)` initializes the sum at `1`, while `np.sum(x, 1)` sums along axis `1`), and `np.sum` is aware of multiple array dimensions.

Minimum and Maximum

Similarly, **Python** has **built-in** `min` and `max` functions, used to find the minimum value and maximum value of any given array:

```
In [ ]: min(big_array), max(big_array)
```

```
Out[ ]: (2.0114398036064074e-07, 0.9999997912802653)
```

NumPy's corresponding functions have **similar syntax**, and again operate **much more quickly**:

```
In [ ]: np.min(big_array), np.max(big_array)
```

```
Out[ ]: (2.0114398036064074e-07, 0.9999997912802653)
```

```
In [ ]: %timeit min(big_array)
        %timeit np.min(big_array)
```

72 ms \pm 177 μ s per loop (mean \pm std. dev. of 7 runs, 10 loops each)

564 μ s \pm 3.11 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

For **min**, **max**, **sum**, and **several other NumPy aggregates**, a **shorter syntax** is to use methods of the array object itself:

```
In [ ]: print(big_array.min(), big_array.max(), big_array.sum())
```

2.0114398036064074e-07 0.9999997912802653 499854.0273321711

Whenever possible, **make sure that you are using the NumPy version** of these aggregates when operating on **NumPy arrays**!

Multidimensional Aggregates

One common **type** of aggregation operation is an **aggregate** along a **row** or **column**.

Example:

```
In [9]: M = rng.integers(0, 10, (3, 4))  
print(M)
```

```
[[5 1 0 7]  
 [3 3 2 9]  
 [1 5 2 0]]
```

NumPy **aggregations** will apply across **all elements** of a multidimensional array:

```
In [ ]: M.sum()
```

```
Out[ ]: 45
```

Aggregation functions take an **additional argument** specifying the **axis** along which the aggregate is computed.

For example, we can find the **minimum value** within **each column** by specifying **axis=0** :

```
In [11]: M.min(axis=0)
```

```
Out[11]: array([1, 1, 0, 0], dtype=int64)
```

The function returns **four values**, corresponding to the **four columns of numbers**.

Similarly, we can find the **maximum value** within each row:

```
In [13]: M.max(axis=1)
```

```
Out[13]: array([7, 9, 5], dtype=int64)
```


The way the **axis is specified** here can be **confusing** to users coming from other languages.

The **axis** keyword specifies the **dimension** of the array that will be **collapsed**, rather than the dimension that will be returned.

So, specifying **axis=0** means that **axis 0 will be collapsed**: for two-dimensional arrays, values within each column will be aggregated.

Other Aggregation Functions

NumPy provides several **other aggregation functions** with a **similar API**.

Additionally **most of them** have a **NaN -safe counterpart** that computes the result while ignoring missing values, which are

marked by the special IEEE floating-point `NaN` value (see **Handling Missing Data**).

The following table provides a **list of useful aggregation functions** available in **NumPy**:

Function name	NaN-safe version	Description
<code>np.sum</code>	<code>np.nansum</code>	Compute sum of elements
<code>np.prod</code>	<code>np.nanprod</code>	Compute product of elements
<code>np.mean</code>	<code>np.nanmean</code>	Compute mean of elements
<code>np.std</code>	<code>np.nanstd</code>	Compute standard deviation
<code>np.var</code>	<code>np.nanvar</code>	Compute variance

Function name	NaN-safe version	Description
<code>np.min</code>	<code>np.nanmin</code>	Find minimum value
<code>np.max</code>	<code>np.nanmax</code>	Find maximum value
<code>np.argmin</code>	<code>np.nanargmin</code>	Find index of minimum value
<code>np.argmax</code>	<code>np.nanargmax</code>	Find index of maximum value
<code>np.median</code>	<code>np.nanmedian</code>	Compute median of elements
<code>np.percentile</code>	<code>np.nanpercentile</code>	Compute rank-based statistics of elements
<code>np.any</code>	N/A	Evaluate whether any elements are true

Function name	NaN-safe version	Description
<code>np.all</code>	N/A	Evaluate whether all elements are true

You will see these aggregates often throughout the rest of the book.

Example: What Is the Average Height of US Presidents?

Aggregates available in NumPy can act as **summary statistics** for a set of values.

As a small example, let's consider **the heights of all US presidents**.

This data is available in the file **president_heights.csv**, which is a **comma-separated list** of labels and values:

```
In [ ]: !head -4 data/president_heights.csv
```

```
order,name,height(cm)
1,George Washington,189
2,John Adams,170
3,Thomas Jefferson,189
```

We'll use the **Pandas package**, which we'll explore more fully in **Part 3**, to **read** the file and **extract this information** (note that the heights are measured in centimeters):

```
In [ ]: import pandas as pd
data = pd.read_csv('data/president_heights.csv')
heights = np.array(data['height(cm)'])
print(heights)
```

```
[189 170 189 163 183 171 185 168 173 183 173 173 175 178 183 19
3 178 173
 174 183 183 168 170 178 182 180 183 178 182 188 175 179 183 19
3 182 183
 177 185 188 188 182 185 191 182]
```

Now that we have this data array, we can **compute a variety of summary statistics**:

```
In [ ]: print("Mean height:      ", heights.mean())
        print("Standard deviation:", heights.std())
        print("Minimum height:   ", heights.min())
        print("Maximum height:   ", heights.max())
```

```
Mean height:      180.04545454545453
Standard deviation: 6.983599441335736
Minimum height:   163
Maximum height:   193
```

Note that in each case, the **aggregation** operation **reduced the entire array to a single summarizing value**, which gives us

information about the **distribution** of values.

We may also wish to compute **quantiles**:

```
In [ ]: print("25th percentile: ", np.percentile(heights, 25))  
        print("Median:         ", np.median(heights))  
        print("75th percentile: ", np.percentile(heights, 75))
```

```
25th percentile:    174.75  
Median:             182.0  
75th percentile:    183.5
```

We see that the **median height of US presidents is 182 cm**, or just shy of six feet.

Of course, sometimes it's **more useful** to see a **visual representation** of this data.

Which we can accomplish using tools in **Matplotlib** (we'll discuss Matplotlib more fully in **Part 4**).

For example, this code generates the following chart:

```
In [ ]: %matplotlib inline
import matplotlib.pyplot as plt
plt.style.use('seaborn-whitegrid')
```

```
In [ ]: plt.hist(heights)
plt.title('Height Distribution of US Presidents')
plt.xlabel('height (cm)')
plt.ylabel('number');
```


Height Distribution of US Presidents

