

The Basics of NumPy Arrays

Data manipulation in Python is nearly **synonymous** with **NumPy array manipulation**.

Tools like **Pandas** ([Part 3](#)) are built **around the NumPy array**.

We'll cover a few categories of **basic array manipulations** here:

- **Attributes of arrays:** Determining the size, shape, memory consumption, and data types of arrays
- **Indexing of arrays:** Getting and setting the values of individual array elements
- **Slicing of arrays:** Getting and setting smaller subarrays within a larger array
- **Reshaping of arrays:** Changing the shape of a given array

- **Joining and splitting of arrays:** Combining multiple arrays into one, and splitting one array into many

NumPy Array Attributes

First let's discuss some useful **array attributes**.

We'll start by defining **random arrays** of **one, two, and three dimensions**.

We'll use **NumPy's random number generator**,

which we will **seed** with a value in order to ensure that the **same random arrays** are generated **each time this code is run**:

```
In [4]: import numpy as np
rng = np.random.default_rng(seed=1701)  # seed for reproducibility
x1 = rng.integers(10, size=6)  # one-dimensional array
```

```
x2 = rng.integers(10, size=(3, 4)) # two-dimensional array  
x3 = rng.integers(10, size=(3, 4, 5)) # three-dimensional array
```

```
In [7]: x1
```

```
Out[7]: array([9, 4, 0, 3, 8, 6], dtype=int64)
```

```
In [9]: x2
```

```
Out[9]: array([[3, 1, 3, 7],  
               [4, 0, 2, 3],  
               [0, 0, 6, 9]], dtype=int64)
```

```
In [11]: x3
```

```
Out[11]: array([[[4, 3, 5, 5, 0],
                  [8, 3, 5, 2, 2],
                  [1, 8, 8, 5, 3],
                  [0, 0, 8, 5, 8]],

                [[5, 1, 6, 2, 3],
                  [1, 2, 5, 6, 2],
                  [5, 2, 7, 9, 3],
                  [5, 6, 0, 2, 0]],

                [[2, 9, 4, 3, 9],
                  [9, 2, 2, 4, 0],
                  [0, 3, 0, 0, 2],
                  [3, 2, 7, 4, 7]]], dtype=int64)
```

Each array has attributes including

- `ndim` (the number of dimensions),
- `shape` (the size of each dimension),
- `size` (the total size of the array), and
- `dtype` (the type of each element):

```
In [ ]: print("x3 ndim: ", x3.ndim)
        print("x3 shape:", x3.shape)
        print("x3 size: ", x3.size)
        print("dtype:   ", x3.dtype)
```

```
x3 ndim:  3
x3 shape: (3, 4, 5)
x3 size:  60
dtype:    int64
```

Array Indexing: Accessing Single Elements

If you are familiar with **Python's standard list indexing**, indexing in **NumPy** will feel quite **familiar**.

In a one-dimensional array, the **ith value** (counting from zero) can be accessed by specifying the **desired index in square brackets**, just as with Python lists:

```
In [23]: x1
```

```
Out[23]: array([9, 4, 0, 3, 8, 6], dtype=int64)
```

```
In [25]: x1[0]
```

```
Out[25]: 9
```

```
In [27]: x1[4]
```

```
Out[27]: 8
```

To index from the **end of the array**, you can use **negative indices**:

```
In [30]: x1[-1]
```

```
Out[30]: 6
```

```
In [ ]: x1[-2]
```

```
Out[ ]: 8
```

In a **multidimensional array**, items can be accessed using a **comma-separated (row, column) tuple**:

```
In [ ]: x2
```

```
Out[ ]: array([[3, 1, 3, 7],  
              [4, 0, 2, 3],  
              [0, 0, 6, 9]])
```

```
In [ ]: x2[0, 0]
```

```
Out[ ]: 3
```

```
In [ ]: x2[2, 0]
```

```
Out[ ]: 0
```

```
In [ ]: x2[2, -1]
```

```
Out[ ]: 9
```

Values can also be **modified** using any of the **index notation**:

```
In [ ]: x2[0, 0] = 12  
x2
```

```
Out[ ]: array([[12,  1,  3,  7],  
               [ 4,  0,  2,  3],  
               [ 0,  0,  6,  9]])
```

Keep in mind that, **unlike Python lists**, NumPy **arrays** have a **fixed type**.

This means, **for example**, that if you **attempt to insert a floating-point** value into an **integer array**, the value will be **silently truncated**.

```
In [ ]: x1[0] = 3.14159 # this will be truncated!  
x1
```

```
Out[ ]: array([3, 4, 0, 3, 8, 6])
```


Array Slicing: Accessing Subarrays

Just as we can use square brackets to **access individual array elements**, we can also use them to **access subarrays** with the **slice** notation, marked by the **colon (:) character**.

The NumPy slicing syntax **follows that of the standard Python list**; to access a slice of an array `x`, use this:

```
x[start:stop:step]
```

If any of these are **unspecified**, they **default** to the values `start=0`, `stop=<size of dimension>`, `step=1`.

One-Dimensional Subarrays

Here are **some examples** of **accessing elements** in one-dimensional subarrays:

```
In [ ]: x1
```

```
Out[ ]: array([3, 4, 0, 3, 8, 6])
```

```
In [ ]: x1[:3] # first three elements
```

```
Out[ ]: array([3, 4, 0])
```

```
In [ ]: x1[3:] # elements after index 3
```

```
Out[ ]: array([3, 8, 6])
```

```
In [ ]: x1[1:4] # middle subarray
```

```
Out[ ]: array([4, 0, 3])
```

```
In [ ]: x1[::2] # every second element
```

```
Out[ ]: array([3, 0, 8])
```

```
In [ ]: x1[1::2] # every second element, starting at index 1
```

```
Out[ ]: array([4, 3, 6])
```

A potentially **confusing** case is when the **step** value is negative.

In this case, the **defaults** for **start** and **stop** are swapped.

This becomes a **convenient way to reverse an array**:

```
In [ ]: x1
```

```
Out[ ]: array([9, 4, 0, 3, 8, 6])
```

```
In [49]: x1[::-1] # all elements, reversed
```

```
Out[49]: array([6, 8, 3, 0, 4, 9], dtype=int64)
```

```
In [ ]: x1[4::-2] # every second element from index 4, reversed
```

```
Out[ ]: array([8, 0, 9])
```

```
In [70]: x1
```

```
Out[70]: array([9, 4, 0, 3, 8, 6], dtype=int64)
```

```
In [68]: x1[4:2:-1] # no defaults here
```

```
Out[68]: array([8, 3], dtype=int64)
```

Multidimensional Subarrays

Multidimensional slices work in the **same way**, with multiple slices separated by **commas**.

For example:

```
In [ ]: x2
```

```
Out[ ]: array([[3, 1, 3, 7],  
               [4, 0, 2, 3],  
               [0, 0, 6, 9]])
```

```
In [ ]: x2[:2, :3] # first two rows & three columns
```

```
Out[ ]: array([[3, 1, 3],  
               [4, 0, 2]])
```

```
In [ ]: x2[:3, ::2] # three rows, every second column
```

```
Out[ ]: array([[3, 3],  
               [4, 2],  
               [0, 6]])
```

```
In [ ]: x2[::-1, ::-1] # all rows & columns, reversed
```

```
Out[ ]: array([[9, 6, 0, 0],  
               [3, 2, 0, 4],  
               [7, 3, 1, 3]])
```

Accessing array rows and columns

One **commonly needed routine** is **accessing single rows or columns** of an array.

This can be done by **combining indexing and slicing**, using an **empty slice** marked by a **single colon (:)**:

```
In [ ]: x2
```

```
Out[ ]: array([[3, 1, 3, 7],  
              [4, 0, 2, 3],  
              [0, 0, 6, 9]])
```

```
In [ ]: x2[:, 0] # first column of x2
```

```
Out[ ]: array([3, 4, 0])
```

```
In [ ]: x2[0, :] # first row of x2
```

```
Out[ ]: array([3, 1, 3, 7])
```

In the case of row access, the **empty slice can be omitted** for a more **compact syntax**:

```
In [ ]: x2[0] # equivalent to x2[0, :]
```

```
Out[ ]: array([12,  1,  3,  7])
```

Subarrays as No-Copy Views

Unlike Python list slices, **NumPy array slices** are returned as **views** rather than **copies** of the array data.

Consider our two-dimensional array from before:

```
In [74]: print(x2)
```

```
[[3 1 3 7]
 [4 0 2 3]
 [0 0 6 9]]
```

Let's extract a 2×2 subarray from this:

```
In [76]: x2_sub = x2[:2, :2]
         print(x2_sub)
```

```
[[3 1]
 [4 0]]
```

Now if we **modify this subarray**, we'll see that **the original array is changed!** Observe:

```
In [78]: x2_sub[0, 0] = 99
         print(x2_sub)
```

```
[[99  1]
 [ 4  0]]
```

```
In [80]: print(x2)
```

```
[[99  1  3  7]
 [ 4  0  2  3]
 [ 0  0  6  9]]
```


Some users may find this surprising, but it **can be advantageous**.

For example, when working with **large datasets**, we can **access and process pieces** of these datasets **without the need to copy** the underlying data buffer.

Creating Copies of Arrays

Despite the nice features of array views, it is sometimes useful to instead **explicitly copy the data** within an array or a subarray.

This can be most easily done with the **copy method**:

```
In [82]: x2_sub_copy = x2[:2, :2].copy()  
         print(x2_sub_copy)
```

```
[[99  1]  
 [ 4  0]]
```

If we now **modify this subarray**, the **original array** is **not touched**:

```
In [84]: x2_sub_copy[0, 0] = 42  
         print(x2_sub_copy)
```

```
[[42  1]  
 [ 4  0]]
```

```
In [86]: print(x2)
```

```
[[99  1  3  7]  
 [ 4  0  2  3]  
 [ 0  0  6  9]]
```

Reshaping of Arrays

Another useful type of operation is **reshaping of arrays**, which can be done with the **reshape method**.

For **example**, if you want to put the numbers 1 through 9 in a 3×3 grid, you can do the following:

```
In [ ]: np.arange(1, 10)
```

```
Out[ ]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [ ]: grid = np.arange(1, 10).reshape(3, 3)  
print(grid)
```

```
[[1 2 3]  
 [4 5 6]  
 [7 8 9]]
```

Note that for this **to work**, the **size of the initial array** must **match** the **size of the reshaped array**.

In most cases the **reshape** method will **return a no-copy view** of the initial array.

A **common reshaping operation** is **converting** a **one-dimensional array** into a **two-dimensional row or column matrix**:

```
In [ ]: x = np.array([1, 2, 3])  
x.reshape((1, 3)) # row vector via reshape
```

```
Out[ ]: array([[1, 2, 3]])
```

```
In [ ]: x.reshape((3, 1)) # column vector via reshape
```

```
Out[ ]: array([[1],  
              [2],  
              [3]])
```

A **convenient shorthand** for this is to use `np.newaxis` in the **slicing syntax**:

```
In [ ]: x[np.newaxis, :] # row vector via newaxis
```

```
Out[ ]: array([[1, 2, 3]])
```

```
In [ ]: x[:, np.newaxis] # column vector via newaxis
```

```
Out[ ]: array([[1],  
              [2],  
              [3]])
```

This is a pattern that we will **utilize** often throughout the **remainder of the book**.

Array Concatenation and Splitting

All of the **preceding routines** worked on **single arrays**.

NumPy also provides tools to **combine multiple arrays** into one, and to conversely **split a single array** into multiple arrays.

Concatenation of Arrays

Concatenation, or joining of two arrays in NumPy, is primarily accomplished using the routines `np.concatenate`, `np.vstack`, and `np.hstack`.

`np.concatenate` takes a tuple or **list of arrays** as its **first argument**, as you can see here:

```
In [91]: x = np.array([1, 2, 3])  
y = np.array([3, 2, 1])  
np.concatenate([x, y])
```

```
Out[91]: array([1, 2, 3, 3, 2, 1])
```

You can also **concatenate more than two arrays** at once:

```
In [93]: z = np.array([99, 99, 99])  
print(np.concatenate([x, y, z]))
```

```
[ 1  2  3  3  2  1 99 99 99]
```

And it can be used for **two-dimensional arrays**:

```
In [95]: grid = np.array([[1, 2, 3],  
                           [4, 5, 6]])
```

```
In [97]: # concatenate along the first axis  
np.concatenate([grid, grid])
```

```
Out[97]: array([[1, 2, 3],  
                [4, 5, 6],  
                [1, 2, 3],  
                [4, 5, 6]])
```

```
In [99]: # concatenate along the second axis (zero-indexed)  
np.concatenate([grid, grid], axis=1)
```

```
Out[99]: array([[1, 2, 3, 1, 2, 3],  
                [4, 5, 6, 4, 5, 6]])
```

For working with **arrays of mixed dimensions**, it can be **clearer** to use the **np.vstack** (vertical stack) and **np.hstack** (horizontal stack) functions:

```
In [101... x
```

```
Out[101...  array([1, 2, 3])
```

```
In [107...  grid
```

```
Out[107...  array([[1, 2, 3],  
          [4, 5, 6]])
```

```
In [111...  # vertically stack the arrays  
np.vstack([x, grid])
```

```
Out[111...  array([[1, 2, 3],  
          [1, 2, 3],  
          [4, 5, 6]])
```

```
In [115...  # horizontally stack the arrays  
y = np.array([[99],  
              [99]])  
np.hstack([grid, y])
```

```
Out[115...  array([[ 1,  2,  3, 99],  
          [ 4,  5,  6, 99]])
```


Splitting of Arrays

The opposite of concatenation is **splitting**, which is implemented by the functions `np.split`, `np.hsplit`, and `np.vsplit`.

For each of these, we can **pass a list of indices** giving the **split points**:

```
In [ ]: x = [1, 2, 3, 99, 99, 3, 2, 1]
        x1, x2, x3 = np.split(x, [3, 5])
        print(x1, x2, x3)
```

```
[1 2 3] [99 99] [3 2 1]
```

Notice that N **split points** leads to $N + 1$ **subarrays**.

The related functions `np.hsplit` and `np.vsplit` are similar:

```
In [121... grid = np.arange(16).reshape((4, 4))
            grid
```

```
Out[121... array([[ 0,  1,  2,  3],  
        [ 4,  5,  6,  7],  
        [ 8,  9, 10, 11],  
        [12, 13, 14, 15]])
```

```
In [123... upper, lower = np.vsplit(grid, [2])  
print(upper)  
print(lower)
```

```
[[0 1 2 3]  
 [4 5 6 7]]  
[[ 8  9 10 11]  
 [12 13 14 15]]
```

```
In [125... left, right = np.hsplit(grid, [2])  
print(left)  
print(right)
```

```
[[ 0  1]
 [ 4  5]
 [ 8  9]
 [12 13]]
[[ 2  3]
 [ 6  7]
 [10 11]
 [14 15]]
```